

Fashion-MNIST Wei Pan

June 15, 2019

Reference for this script: <https://github.com/ageron/handson-ml>
<https://gist.github.com/hitvoice/36cf44689065ca9b927431546381a3f7>/In Fashion MNIST dataset, each training and test example is assigned to one of the following labels:
0 T-shirt/top 1 Trouser 2 Pullover 3 Dress 4 Coat 5 Sandal 6 Shirt 7 Sneaker 8 Bag 9 Ankle boot

0.1 Setup

```
In [1]: # To support both python 2 and python 3
        from __future__ import division, print_function, unicode_literals

        # Common imports
        import numpy as np
        import pandas as pd
        import os
        import warnings
        warnings.filterwarnings("ignore")

        import tensorflow as tf
        from tensorflow import keras
        from tensorflow_graph_in_jupyter import show_graph

        from sklearn.model_selection import cross_val_score
        from sklearn.model_selection import cross_val_predict
        from sklearn.metrics import precision_score, recall_score
        from sklearn.metrics import confusion_matrix

        # to make this notebook's output stable across runs
        np.random.seed(42)

        # To plot pretty figures
        import seaborn as sns
        %matplotlib inline
        import matplotlib
        import matplotlib.pyplot as plt
        plt.rcParams['axes.labelsize'] = 14
        plt.rcParams['xtick.labelsize'] = 12
        plt.rcParams['ytick.labelsize'] = 12
```

```
# Where to save the figures
```

```
PROJECT_ROOT_DIR = "."
```

```
def save_fig(fig_id, tight_layout=True):  
    path = os.path.join(PROJECT_ROOT_DIR, fig_id + ".png")  
    print("Saving figure", fig_id)  
    if tight_layout:  
        plt.tight_layout()  
    plt.savefig(path, format='png', dpi=600)
```

```
In [2]: print(tf.__version__)
```

```
1.13.0-rc2
```

```
In [3]: # To plot images
```

```
def plot_images(instances, images_per_row=10, **options):  
    size = 28  
    images_per_row = min(len(instances), images_per_row)  
    images = [instance.reshape(size,size) for instance in instances]  
    n_rows = (len(instances) - 1) // images_per_row + 1  
    row_images = []  
    n_empty = n_rows * images_per_row - len(instances)  
    images.append(np.zeros((size, size * n_empty)))  
    for row in range(n_rows):  
        rimages = images[row * images_per_row : (row + 1) * images_per_row]  
        row_images.append(np.concatenate(rimages, axis=1))  
    image = np.concatenate(row_images, axis=0)  
    plt.imshow(image, cmap = matplotlib.cm.binary, **options)  
    plt.axis("off")
```

```
In [4]: # To plot confusion matrix
```

```
# Source: https://gist.github.com/hitvoice/36cf44689065ca9b927431546381a3f7/
```

```
def confusion_matrix_plot(y_true, y_pred, filename, labels=range(10), ymap=None, figsize=(10, 10),  
                          midpoint=500, color=None, v_min=0, v_max=1000):  
    if ymap is not None:  
        y_pred = [ymap[yi] for yi in y_pred]  
        y_true = [ymap[yi] for yi in y_true]  
        labels = [ymap[yi] for yi in labels]  
    cm = confusion_matrix(y_true, y_pred, labels=labels)  
    cm_sum = np.sum(cm, axis=1, keepdims=True)  
    cm_perc = cm / cm_sum.astype(float) * 100  
    annot = np.empty_like(cm).astype(str)  
    nrows, ncols = cm.shape  
    for i in range(nrows):  
        for j in range(ncols):  
            c = cm[i, j]  
            p = cm_perc[i, j]  
            if i == j:  
                s = "diag: %d%%" % p
```

```

        s = cm_sum[i]
        annot[i, j] = '%.1f%%\n%d/%d' % (p, c, s)
    elif c == 0:
        annot[i, j] = ''
    else:
        annot[i, j] = '%.1f%%\n%d' % (p, c)
cm = pd.DataFrame(cm, index=labels, columns=labels)
cm.index.name = 'Actual'
cm.columns.name = 'Predicted'
fig, ax = plt.subplots(figsize=figsize)
sns.heatmap(cm, annot=annot, fmt='', ax=ax, annot_kws={"size": 13},
            cmap=color, center=midpoint, vmin=v_min, vmax=v_max)
plt.savefig(filename)

```

0.2 Get the Data

```

In [5]: # import data
        zalando_data_train = pd.read_csv('fashion-mnist_train.csv')
        zalando_data_test = pd.read_csv('fashion-mnist_test.csv')

In [6]: y_train_raw = zalando_data_train['label'].values
        print(y_train_raw.shape)
        X_train_raw = zalando_data_train.drop('label', axis=1).values
        print(X_train_raw.shape)
        y_test_raw = zalando_data_test['label'].values
        print(y_test_raw.shape)
        X_test_raw = zalando_data_test.drop('label', axis=1).values
        print(X_test_raw.shape)

(60000,)
(60000, 784)
(10000,)
(10000, 784)

```

```

In [7]: print (zalando_data_train['label'].value_counts())

```

```

9    6000
8    6000
7    6000
6    6000
5    6000
4    6000
3    6000
2    6000
1    6000
0    6000
Name: label, dtype: int64

```

```
In [8]: # show some images
plt.figure(figsize=(9,9))
example_images = np.r_[X_train_raw[:12000:600], X_train_raw[13000:30600:600], X_train_
plot_images(example_images, images_per_row=10)
save_fig("images_plot")
plt.show()
```

Saving figure images_plot



```
In [9]: # plot images for each label
X_0 = X_train_raw[(y_train_raw == 0)]
X_1 = X_train_raw[(y_train_raw == 1)]
X_2 = X_train_raw[(y_train_raw == 2)]
```


```

X_3 = X_train_raw[(y_train_raw == 3)]
X_4 = X_train_raw[(y_train_raw == 4)]
X_5 = X_train_raw[(y_train_raw == 5)]
X_6 = X_train_raw[(y_train_raw == 6)]
X_7 = X_train_raw[(y_train_raw == 7)]
X_8 = X_train_raw[(y_train_raw == 8)]
X_9 = X_train_raw[(y_train_raw == 9)]

plt.figure(figsize=(24,24))
plt.subplot(521); plot_images(X_0[:25], images_per_row=5)
plt.subplot(522); plot_images(X_1[:25], images_per_row=5)
plt.subplot(523); plot_images(X_2[:25], images_per_row=5)
plt.subplot(524); plot_images(X_3[:25], images_per_row=5)
plt.subplot(525); plot_images(X_4[:25], images_per_row=5)
plt.subplot(526); plot_images(X_5[:25], images_per_row=5)
plt.subplot(527); plot_images(X_6[:25], images_per_row=5)
plt.subplot(528); plot_images(X_7[:25], images_per_row=5)
plt.subplot(529); plot_images(X_8[:25], images_per_row=5)
plt.subplot(5,2,10); plot_images(X_9[:25], images_per_row=5)
save_fig("images_for_each_label")
plt.show()

```

Saving figure images_for_each_label



0.3 Linear Classifier

```
In [10]: # prepare data
        y_train = y_train_raw
        X_train = X_train_raw
        y_test = y_test_raw
        X_test = X_test_raw

In [11]: # prepare Logistic Regression classifier
        from sklearn.linear_model import LogisticRegression
        logit_clf = LogisticRegression(tol=0.0001, C=1.0, random_state=42, solver='sag',
                                       max_iter=300, multi_class='multinomial')


In [12]: # training and validation
        logit_clf.fit(X_train, y_train)
        y_train_logit_pred = logit_clf.predict(X_train)
        y_test_logit_pred = logit_clf.predict(X_test)

        print("Training Accuracy", (y_train_logit_pred == y_train).mean())
        print("Test Accuracy", (y_test_logit_pred == y_test).mean())

/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/sag.py:334: ConvergenceWarning: The
"the coef_ did not converge", ConvergenceWarning)
```

Training Accuracy 0.8794166666666666
Test Accuracy 0.8462

```
In [13]: # Evaluating classifier
        confusion_matrix_plot(y_test, y_test_logit_pred,
                              filename="confusion_matrix_logit_test",
                              ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover", 3:"Dress", 4:"Coat", 5:"Sandal",
                                    6:"Shirt", 7:"Sneaker", 8:"Bag", 9:"Ankle boot"},
                              color=sns.diverging_palette(20, 220, sep=10, n=25, s=90, l=45, center='light'),
                              midpoint=75, figsize=(15, 12), v_min=0, v_max=1000)
```



```
In [14]: # Linear SVM classifier
from sklearn.linear_model import SGDClassifier
sgd_svm_clf = SGDClassifier(max_iter=150, random_state=42, loss='hinge', tol=0.01)
sgd_svm_clf.fit(X_train, y_train)
```

/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/stochastic_gradient.py:603: ConvergenceWarning)

```
Out[14]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
    l1_ratio=0.15, learning_rate='optimal', loss='hinge', max_iter=150,
    n_iter=None, n_iter_no_change=5, n_jobs=None, penalty='l2',
    power_t=0.5, random_state=42, shuffle=True, tol=0.01,
    validation_fraction=0.1, verbose=0, warm_start=False)
```

```
In [15]: # training and validation
y_train_sgd_svm_pred = sgd_svm_clf.predict(X_train)
```



```

y_test_sgd_svm_pred = sgd_svm_clf.predict(X_test)

print("Training Accuracy", (y_train_sgd_svm_pred == y_train).mean())
print("Test Accuracy", (y_test_sgd_svm_pred == y_test).mean())

```

Training Accuracy 0.8507833333333333


Test Accuracy 0.8304

In [16]: # Evaluating classifier

```

confusion_matrix_plot(y_test, y_test_sgd_svm_pred,
    filename="confusion_matrix_sgd_svm_test",
    ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover", 3:"Dress", 4:"Coat", 5:"Sandal",
    6:"Shirt", 7:"Sneaker", 8:"Bag", 9:"Ankle boot"},
    color=sns.diverging_palette(20,220, sep=10, n=25, s=90, l=45, center='light'),
    midpoint=75, figsize=(15,12), v_min=0, v_max=1000)

```



0.3.1 Linear SVM Classifier (pre-processing: scaling)

```
In [17]: # prepare classifier
        from sklearn.preprocessing import StandardScaler
        scaler = StandardScaler()
        X_train_scaled = scaler.fit_transform(X_train.astype(np.float64))
        X_test_scaled = scaler.fit_transform(X_test.astype(np.float64))

In [18]: # training and validation
        sgd_svm_clf = SGDClassifier(max_iter=150, random_state=42, loss='hinge', tol=0.0001)
        sgd_svm_clf.fit(X_train_scaled, y_train)
        y_train_svm_scaled_pred = sgd_svm_clf.predict(X_train_scaled)
        y_test_svm_scaled_pred = sgd_svm_clf.predict(X_test_scaled)


        print("Training Accuracy", (y_train_svm_scaled_pred == y_train).mean())
        print("Test Accuracy", (y_test_svm_scaled_pred == y_test).mean())
```

Training Accuracy 0.8434833333333334

Test Accuracy 0.8361

/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/stochastic_gradient.py:603: ConvergenceWarning

```
In [19]: # Evaluating classifier
        confusion_matrix_plot(y_test, y_test_svm_scaled_pred,
                              filename="confusion_matrix_sgd_scaled_test",
                              ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover", 3:"Dress", 4:"Coat", 5:"Sandal",
                                     6:"Shirt", 7:"Sneaker", 8:"Bag", 9:"Ankle boot"},
                              color=sns.diverging_palette(20, 220, sep=10, n=25, s=90, l=45, center='light'),
                              midpoint=75, figsize=(15, 12), v_min=0, v_max=1000)
```



0.4 Random Forest Classifier

In [42]: *# Data preparation*

```
y_train = y_train_raw
X_train = X_train_raw
y_test = y_test_raw
X_test = X_test_raw
```

In [21]: *# Classifier preparation*

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.multiclass import OneVsOneClassifier
```

```
# ovo_clf = OneVsOneClassifier(RandomForestClassifier(random_state=42)) # overfitting
ovo_rf_clf = OneVsOneClassifier(RandomForestClassifier(random_state=42, max_depth = 20))
```

In [22]: *# Training and validation*

```
ovo_rf_clf.fit(X_train, y_train)
y_train_rf_pred = ovo_rf_clf.predict(X_train)
```

```

y_test_rf_pred = ovo_rf_clf.predict(X_test)

print("Training Accuracy", (y_train_rf_pred == y_train).mean())
print("Test Accuracy", (y_test_rf_pred == y_test).mean())

```

Training Accuracy 0.9906666666666667


Test Accuracy 0.8631

In [23]: *# Evaluating classifier*

```

confusion_matrix_plot(y_test, y_test_rf_pred,
    filename="confusion_matrix_rf_test",
    ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover", 3:"Dress", 4:"Coat", 5:"Sandal",
        6:"Shirt", 7:"Sneaker", 8:"Bag", 9:"Ankle boot"},
    color=sns.diverging_palette(20,220, sep=10, n=25, s=90, l=45, center='light'),
    midpoint=75, figsize=(15,12), v_min=0, v_max=1000)

```



0.5 CNN

```
In [25]: def reset_graph(seed=42):
         tf.reset_default_graph()
         tf.set_random_seed(seed)
         np.random.seed(seed)

In [26]: from datetime import datetime

         now = datetime.utcnow().strftime("%Y%m%d%H%M%S")
         root_logdir = "."
         logdir = "{}//run-{}//".format(root_logdir, now)

In [ ]: # Building CNN architecture
         height = 28
         width = 28
         channels = 1
         n_inputs = height * width

         conv1_fmaps = 64#16
         conv1_ksize = 3
         conv1_stride = 1
         conv1_pad = "SAME"

         conv2_fmaps = 16
         conv2_ksize = 3
         conv2_stride = 2
         conv2_pad = "SAME"

         pool3_fmaps = 16#8

         n_fc1 = 64#32
         n_outputs = 10

         reset_graph()

         with tf.name_scope("inputs"):
             X = tf.placeholder(tf.float32, shape=[None, n_inputs], name="X")
             X_reshaped = tf.reshape(X, shape=[-1, height, width, channels])
             y = tf.placeholder(tf.int32, shape=[None], name="y")

         conv1 = tf.layers.conv2d(X_reshaped, filters=conv1_fmaps, kernel_size=conv1_ksize,
                                   strides=conv1_stride, padding=conv1_pad,
                                   activation=tf.nn.relu, name="conv1")
         conv2 = tf.layers.conv2d(conv1, filters=conv2_fmaps, kernel_size=conv2_ksize,
                                   strides=conv2_stride, padding=conv2_pad,
                                   activation=tf.nn.relu, name="conv2")

         with tf.name_scope("pool3"):
```

```

pool3 = tf.nn.max_pool(conv2, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding="V
pool3_flat = tf.reshape(pool3, shape=[-1, pool3_fmaps * 7 * 7])

with tf.name_scope("fc1"):
    fc1 = tf.layers.dense(pool3_flat, n_fc1, activation=tf.nn.relu, name="fc1")

with tf.name_scope("output"):
    logits = tf.layers.dense(fc1, n_outputs, name="output")
    Y_proba = tf.nn.softmax(logits, name="Y_proba")

with tf.name_scope("train"):
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(logits=logits, labels=y)
    loss = tf.reduce_mean(xentropy)
    optimizer = tf.train.AdamOptimizer()
    training_op = optimizer.minimize(loss)

with tf.name_scope("eval"):
    correct = tf.nn.in_top_k(logits, y, 1)
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

with tf.name_scope("init_and_save"):
    init = tf.global_variables_initializer()
    saver = tf.train.Saver()

file_writer = tf.summary.FileWriter(logdir , tf.get_default_graph()) # for TensorBoard

```

In [28]: *# Data preparation*

```

X_train = X_train_raw.astype(np.float32).reshape(-1, 28*28) / 255.0
X_test = X_test_raw.astype(np.float32).reshape(-1, 28*28) / 255.0
y_train = y_train_raw.astype(np.int32)
y_test = y_test_raw.astype(np.int32)
X_valid, X_train = X_train[:5000], X_train[5000:]
y_valid, y_train = y_train[:5000], y_train[5000:]
X_train.shape

```

Out[28]: (55000, 784)

In [29]: *def shuffle_batch(X, y, batch_size):*

```

    rnd_idx = np.random.permutation(len(X))
    n_batches = len(X) // batch_size
    for batch_idx in np.array_split(rnd_idx, n_batches):
        X_batch, y_batch = X[batch_idx], y[batch_idx]
        yield X_batch, y_batch

```

In [30]: *# Trainning CNN model*

```

n_epochs = 30
batch_size = 100

with tf.Session() as sess:

```

```

init.run()
for epoch in range(n_epochs):
    for X_batch, y_batch in shuffle_batch(X_train, y_train, batch_size):
        sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
        acc_batch = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
        acc_test = accuracy.eval(feed_dict={X: X_test, y: y_test})
        print(epoch, "Last batch accuracy:", acc_batch, "Test accuracy:", acc_test)
    saver.save(sess, "./fashion_mnist_cnn_model.ckpt")

file_writer.close() # for TensorBoard

```

0 Last batch accuracy: 0.77 Test accuracy: 0.8679
1 Last batch accuracy: 0.91 Test accuracy: 0.8842
2 Last batch accuracy: 0.87 Test accuracy: 0.8985
3 Last batch accuracy: 0.9 Test accuracy: 0.8928
4 Last batch accuracy: 0.95 Test accuracy: 0.9063
5 Last batch accuracy: 0.96 Test accuracy: 0.9054
6 Last batch accuracy: 0.95 Test accuracy: 0.9095
7 Last batch accuracy: 0.92 Test accuracy: 0.9142
8 Last batch accuracy: 0.95 Test accuracy: 0.9168
9 Last batch accuracy: 0.93 Test accuracy: 0.9149
10 Last batch accuracy: 0.98 Test accuracy: 0.9166
11 Last batch accuracy: 0.94 Test accuracy: 0.9168
12 Last batch accuracy: 0.98 Test accuracy: 0.9158
13 Last batch accuracy: 0.94 Test accuracy: 0.9126
14 Last batch accuracy: 0.98 Test accuracy: 0.9143
15 Last batch accuracy: 0.97 Test accuracy: 0.9114
16 Last batch accuracy: 0.93 Test accuracy: 0.9174
17 Last batch accuracy: 0.93 Test accuracy: 0.9182
18 Last batch accuracy: 0.93 Test accuracy: 0.9194
19 Last batch accuracy: 0.93 Test accuracy: 0.9173
20 Last batch accuracy: 0.92 Test accuracy: 0.9138
21 Last batch accuracy: 0.98 Test accuracy: 0.9158
22 Last batch accuracy: 0.99 Test accuracy: 0.9118
23 Last batch accuracy: 0.95 Test accuracy: 0.9164
24 Last batch accuracy: 0.97 Test accuracy: 0.9113
25 Last batch accuracy: 0.94 Test accuracy: 0.9068
26 Last batch accuracy: 0.95 Test accuracy: 0.9141
27 Last batch accuracy: 0.95 Test accuracy: 0.9123
28 Last batch accuracy: 0.98 Test accuracy: 0.9144
29 Last batch accuracy: 1.0 Test accuracy: 0.9133

```

In [31]: from tensorflow_graph_in_jupyter import show_graph
         show_graph(tf.get_default_graph())

```


<IPython.core.display.HTML object>

```
In [32]: # Validation on test set
with tf.Session() as sess:
    saver.restore(sess, "./fashion_mnist_cnn_model.ckpt")
    Z = logits.eval(feed_dict={X: X_test})
    y_test_cnn_pred = np.argmax(Z, axis=1)

    print("Test Accuracy", (y_test_cnn_pred == y_test).mean())# accuracy for test set
```

WARNING:tensorflow:From /anaconda3/lib/python3.6/site-packages/tensorflow/python/training/saver.py:1268: tf.nn.conv2d is deprecated and will be removed in a future version. Instructions for updating:
Use standard file APIs to check for files with this prefix.
INFO:tensorflow:Restoring parameters from ./fashion_mnist_cnn_model.ckpt
Test Accuracy 0.9133

```
In [33]: # Evaluating classifier
confusion_matrix_plot(y_test, y_test_cnn_pred,
    filename="confusion_matrix_cnn_test",
    ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover",3:"Dress",4:"Coat",5:"Sandal",
    6:"Shirt",7:"Sneaker",8:"Bag",9:"Ankle boot"},
    color=sns.diverging_palette(20,220, sep=10, n=25, s=90, l=45,center='light'),
    midpoint=75,figsize=(15,12),v_min=0,v_max=1000)
```

0.6 Neural Nets

The training time is too long and results we obtain is not that good.

```
In [61]: def reset_graph(seed=42):
         tf.reset_default_graph()
         tf.set_random_seed(seed)
         np.random.seed(seed)
```

```
In [62]: from datetime import datetime

         now = datetime.utcnow().strftime("%Y%m%d%H%M%S")
         root_logdir = "."
         logdir = "{}//run-{}//".format(root_logdir, now)
```

```
In [63]: #define validation and training sets
         X_test = X_test.astype(np.int32)
         y_train = y_train.astype(np.int32)
```

```

y_test = y_test.astype(np.int32)
X_valid, X_train = X_train[:100], X_train[100:]
y_valid, y_train = y_train[:100], y_train[100:]

n_inputs = 784 # number of features
n_hidden1 = 147
n_hidden2 = 70
# dedfine more hidden layers
n_hidden3 = 50
n_outputs = 10
tf.reset_default_graph() # reset graph

# y place holder for output
# auxialrry variable
X = tf.placeholder(tf.float32, shape=(None, n_inputs), name="X")
y = tf.placeholder(tf.int32, shape=(None), name="y")
a = tf.placeholder(tf.float32, shape=(None), name='a')

summary1 = tf.reduce_mean(a, name="summary1")

#layer of network takes input of neurones
def neuron_layer(X, n_neurons, name, activation=None):
    with tf.name_scope(name):
        n_inputs = int(X.get_shape()[1])
        stddev = 2 / np.sqrt(n_inputs)
        init = tf.truncated_normal((n_inputs, n_neurons), stddev=stddev)
        W = tf.Variable(init, name="kernel")
        b = tf.Variable(tf.zeros([n_neurons]), name="bias")
        Z = tf.matmul(X, W) + b
        if activation is not None:
            return activation(Z)
        else:
            return Z

with tf.name_scope("dnn"):
    hidden1 = neuron_layer(X, n_hidden1, name="hidden1", activation=tf.nn.relu)
    hidden2 = neuron_layer(hidden1, n_hidden2, name="hidden2", activation=tf.nn.relu)
    hidden3 = neuron_layer(hidden2, n_hidden3, name="hidden3", activation=tf.nn.relu)
    logits = neuron_layer(hidden3, n_outputs, name="outputs")

with tf.name_scope("loss"):
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y, logits=logits)
    loss = tf.reduce_mean(xentropy, name="loss")

```

```

#learning rate - change if issues regrading convergence or others happen
learning_rate = 1e-1
with tf.name_scope("train"):
    optimizer = tf.train.AdagradOptimizer(learning_rate)
    training_op = optimizer.minimize(loss)

#define accuracy
with tf.name_scope("accuracy"):
    correct = tf.nn.in_top_k(logits, y, 1)
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

#leave this
def shuffle_batch(X, y, batch_size):
    rnd_idx = np.random.permutation(len(X))
    n_batches = len(X) // batch_size
    for batch_idx in np.array_split(rnd_idx, n_batches):
        X_batch, y_batch = X[batch_idx], y[batch_idx]
        yield X_batch, y_batch

```

```

In [ ]: from random import randint
        n_epochs = 100
        batch_size = 100

```

```

init = tf.global_variables_initializer()
saver = tf.train.Saver()

```

```

summary2 = tf.summary.scalar('Accuracy', summary1)
file_writer = tf.summary.FileWriter(logdir, tf.get_default_graph())

```

```

with tf.Session() as sess:
    init.run()
    acc_val = accuracy.eval(feed_dict={X: X_valid, y: y_valid})
    print("Val accuracy init:", acc_val)
    for epoch in range(n_epochs):
        for X_batch, y_batch in shuffle_batch(X_train, y_train, batch_size):
            sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
            acc_batch = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
            acc_val = accuracy.eval(feed_dict={X: X_valid, y: y_valid})
            print(epoch, "Batch accuracy:", acc_batch, "Val accuracy:", acc_val)

        summary_str = summary2.eval(feed_dict={a: acc_val })
        file_writer.add_summary(summary_str, epoch)

    saver.save(sess, "./fashion_mnist_dnn_model.ckpt")

```

```

In [65]: with tf.Session() as sess:
        saver.restore(sess, "./fashion_mnist_dnn_model.ckpt")


```

```
Z = logits.eval(feed_dict={X: X_test})
y_test_dnn_pred = np.argmax(Z, axis=1)
```

INFO:tensorflow:Restoring parameters from ./fashion_mnist_dnn_model.ckpt

In [66]: # *Evaluating classifier*

```
confusion_matrix_plot(y_test, y_test_dnn_pred,
    filename="confusion_matrix_dnn_test",
    ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover",3:"Dress",4:"Coat",5:"Sandal",
    6:"Shirt",7:"Sneaker",8:"Bag",9:"Ankle boot"},
    color=sns.diverging_palette(20,220, sep=10, n=25, s=90, l=45,center = 'light'),
    midpoint=75,figsize=(15,12),v_min=0,v_max=1000)
```



0.7 RNN

The training time is too long and results we obtain is not that good.

```

In [ ]: import tensorflow as tf

reset_graph()

n_steps = 28
n_inputs = 28
n_neurons = 150
n_outputs = 10

learning_rate = 0.001

X = tf.placeholder(tf.float32, [None, n_steps, n_inputs])
y = tf.placeholder(tf.int32, [None])

basic_cell = tf.contrib.rnn.BasicRNNCell(num_units=n_neurons)
outputs, states = tf.nn.dynamic_rnn(basic_cell, X, dtype=tf.float32)

logits = tf.layers.dense(states, n_outputs)
xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y,
                                                            logits=logits)

prediction = tf.argmax(logits, 1)

loss = tf.reduce_mean(xentropy)
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
training_op = optimizer.minimize(loss)
correct = tf.nn.in_top_k(logits, y, 1) # check if the prediction is correct [True, False]
accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))

init = tf.global_variables_initializer()

In [ ]: from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets("FashionMnist-data")

X_test = mnist.test.images.reshape((-1, n_steps, n_inputs))
y_test = mnist.test.labels

In [69]: import time
n_epochs = 10
batch_size = 150
start_time = time.time()
saver = tf.train.Saver()

with tf.Session() as sess:
    init.run()
    for epoch in range(n_epochs):
        for iteration in range(mnist.train.num_examples // batch_size):
            X_batch, y_batch = mnist.train.next_batch(batch_size)

```

```

        X_batch = X_batch.reshape((-1, n_steps, n_inputs))
        sess.run(training_op, feed_dict={X: X_batch, y: y_batch})
        # print(iteration,end=' ')
        acc_train = accuracy.eval(feed_dict={X: X_batch, y: y_batch})
        acc_test = accuracy.eval(feed_dict={X: X_test, y: y_test})
        print(epoch, "Test accuracy:", acc_test)
        saver.save(sess, "./fashion_mnist_rnn_model.ckpt")

    print("--- %s seconds ---" % (time.time() - start_time))

0 Test accuracy: 0.7962
1 Test accuracy: 0.8149
2 Test accuracy: 0.8262
3 Test accuracy: 0.8335
4 Test accuracy: 0.8429
5 Test accuracy: 0.8333
6 Test accuracy: 0.8426
7 Test accuracy: 0.8494
8 Test accuracy: 0.8477
9 Test accuracy: 0.8537
--- 83.00638699531555 seconds ---


In [70]: with tf.Session() as sess:
        saver.restore(sess, "./fashion_mnist_rnn_model.ckpt")
        prediction = tf.argmax(logits, axis=1)
        y_test_rnn_pred = sess.run(prediction, feed_dict={X: X_test})

        print("Test Accuracy", (y_test_rnn_pred == y_test).mean()) # accuracy for test set

INFO:tensorflow:Restoring parameters from ./fashion_mnist_rnn_model.ckpt
Test Accuracy 0.8537

In [71]: # Evaluating classifier
        confusion_matrix_plot(y_test, y_test_rnn_pred,
                               filename="confusion_matrix_rnn_test",
                               ymap={0:"T-shirt", 1:"Trouser", 2:"Pullover", 3:"Dress", 4:"Coat", 5:"Sandal",
                                     6:"Shirt", 7:"Sneaker", 8:"Bag", 9:"Ankle boot"},
                               color=sns.diverging_palette(20,220, sep=10, n=25, s=90, l=45, center='light'),
                               midpoint=75, figsize=(15,12), v_min=0, v_max=1000)

```



0.8 Error Analysis on training set for Logistic Regression Classifier

In [60]: *# more errors happend to label 0,2,4,6*

```
X_0 = X_train[(y_train == 0)]
X_2 = X_train[(y_train == 2)]
X_4 = X_train[(y_train == 4)]
X_6 = X_train[(y_train == 6)]
```

```
plt.figure(figsize=(7,7))
plt.subplot(221); plot_images(X_0[:25], images_per_row=5)
plt.subplot(222); plot_images(X_2[:25], images_per_row=5)
plt.subplot(223); plot_images(X_4[:25], images_per_row=5)
plt.subplot(224); plot_images(X_6[:25], images_per_row=5)
# save_fig("images_for_labels_0&2&4&6")
```




```
In [59]: # error analysis : label 2 and 4
cl_a, cl_b = 2, 4
X_aa = X_train[(y_train == cl_a) & (y_train_logit_pred == cl_a)]
X_ab = X_train[(y_train == cl_a) & (y_train_logit_pred == cl_b)]
X_ba = X_train[(y_train == cl_b) & (y_train_logit_pred == cl_a)]
X_bb = X_train[(y_train == cl_b) & (y_train_logit_pred == cl_b)]
plt.figure(figsize=(5,7))
plt.subplot(221); plot_images(X_aa[:9], images_per_row=3)
plt.subplot(222); plot_images(X_ab[:9], images_per_row=3)
plt.subplot(223); plot_images(X_ba[:9], images_per_row=3)
plt.subplot(224); plot_images(X_bb[:9], images_per_row=3)
# save_fig("error_analysis_images_plot_2&4")
plt.show()
```




```
In [58]: # error analysis : label 6 and 0
cl_a, cl_b = 6, 0
X_aa = X_train[(y_train == cl_a) & (y_train_logit_pred == cl_a)]
X_ab = X_train[(y_train == cl_a) & (y_train_logit_pred == cl_b)]
X_ba = X_train[(y_train == cl_b) & (y_train_logit_pred == cl_a)]
X_bb = X_train[(y_train == cl_b) & (y_train_logit_pred == cl_b)]

plt.figure(figsize=(7,7))
plt.subplot(221); plot_images(X_aa[:25], images_per_row=5)
plt.subplot(222); plot_images(X_ab[:25], images_per_row=5)
plt.subplot(223); plot_images(X_ba[:25], images_per_row=5)
plt.subplot(224); plot_images(X_bb[:25], images_per_row=5)
# save_fig("error_analysis_images_plot_6&0")
plt.show()
```



```
In [57]: # error analysis : label 6 and 0
cl_a, cl_b = 2, 8
X_aa = X_train[(y_train == cl_a) & (y_train_logit_pred == cl_a)]
X_ab = X_train[(y_train == cl_a) & (y_train_logit_pred == cl_b)]
X_ba = X_train[(y_train == cl_b) & (y_train_logit_pred == cl_a)]
X_bb = X_train[(y_train == cl_b) & (y_train_logit_pred == cl_b)]

plt.figure(figsize=(7,7))
plt.subplot(221); plot_images(X_aa[:16], images_per_row=4)
plt.subplot(222); plot_images(X_ab[:16], images_per_row=4)
plt.subplot(223); plot_images(X_ba[:16], images_per_row=4)
plt.subplot(224); plot_images(X_bb[:16], images_per_row=4)
# save_fig("error_analysis_images_plot_248")
plt.show()
```



0.9 Error Analysis on training set for Logistic Regression Classifier

In [56]: *# more errors happend to label 0,2,4,6*

```
X_0 = X_test[(y_test == 0)]
X_2 = X_test[(y_test == 2)]
X_4 = X_test[(y_test == 4)]
X_6 = X_test[(y_test == 6)]
```

```
plt.figure(figsize=(7,7))
plt.subplot(221); plot_images(X_0[:25], images_per_row=5)
plt.subplot(222); plot_images(X_2[:25], images_per_row=5)
plt.subplot(223); plot_images(X_4[:25], images_per_row=5)
plt.subplot(224); plot_images(X_6[:25], images_per_row=5)
```

```
save_fig("images_for_labels_0&2&4&6")
plt.show()
```

Saving figure images_for_labels_0&2&4&6



In [48]: *# more errors happend to label 5,7,9*

```
X_5 = X_test[(y_test == 5)]
X_7 = X_test[(y_test == 7)]
X_9 = X_test[(y_test == 9)]
```

```
plt.figure(figsize=(9,9))
plt.subplot(131); plot_images(X_5[:25], images_per_row=5)
plt.subplot(132); plot_images(X_7[:25], images_per_row=5)
```

```
plt.subplot(133); plot_images(X_9[:25], images_per_row=5)
save_fig("images_for_labels_5&7&9")
plt.show()
```

Saving figure images_for_labels_5&7&9



```
In [49]: # error analysis : label 5 and 7
cl_a, cl_b = 5, 7
X_aa = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_a)]
X_ab = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_b)]
X_ba = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_a)]
X_bb = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_b)]

plt.figure(figsize=(5,7))
plt.subplot(221); plot_images(X_aa[:9], images_per_row=3)
plt.subplot(222); plot_images(X_ab[:9], images_per_row=3)
plt.subplot(223); plot_images(X_ba[:9], images_per_row=3)
plt.subplot(224); plot_images(X_bb[:9], images_per_row=3)
# save_fig("error_analysis_images_plot_5&7")
plt.show()
```



```
In [53]: # error analysis : label 7 and 9
cl_a, cl_b = 7, 9
X_aa = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_a)]
X_ab = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_b)]
X_ba = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_a)]
X_bb = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_b)]

plt.figure(figsize=(5,7))
plt.subplot(221); plot_images(X_aa[:16], images_per_row=4)
plt.subplot(222); plot_images(X_ab[:16], images_per_row=4)
plt.subplot(223); plot_images(X_ba[:16], images_per_row=4)
plt.subplot(224); plot_images(X_bb[:16], images_per_row=4)
# save_fig("error_analysis_images_plot_7&9")
plt.show()
```



```
In [55]: # error analysis : label 2 and 4
         cl_a, cl_b = 2, 4
         X_aa = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_a)]
         X_ab = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_b)]
         X_ba = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_a)]
         X_bb = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_b)]

         plt.figure(figsize=(5,7))
         plt.subplot(221); plot_images(X_aa[:9], images_per_row=3)
         plt.subplot(222); plot_images(X_ab[:9], images_per_row=3)
         plt.subplot(223); plot_images(X_ba[:9], images_per_row=3)
         plt.subplot(224); plot_images(X_bb[:9], images_per_row=3)
         # save_fig("error_analysis_images_plot_2&4")
         plt.show()
```



```
In [54]: # error analysis : label 6 and 0
         cl_a, cl_b = 6, 0
         X_aa = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_a)]
         X_ab = X_test[(y_test == cl_a) & (y_test_logit_pred == cl_b)]
         X_ba = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_a)]
         X_bb = X_test[(y_test == cl_b) & (y_test_logit_pred == cl_b)]

         plt.figure(figsize=(5,7))
         plt.subplot(221); plot_images(X_aa[:9], images_per_row=3)
         plt.subplot(222); plot_images(X_ab[:9], images_per_row=3)
         plt.subplot(223); plot_images(X_ba[:9], images_per_row=3)
         plt.subplot(224); plot_images(X_bb[:9], images_per_row=3)
         # save_fig("error_analysis_images_plot_6&0")
         plt.show()
```